

k Nearest Neighbour Search with CUDA

Matthias Greshake, Alexander Mock

mgreshake@uos.de, amock@uos.de

Institut für Informatik
AG Wissensbasierte Systeme

26. January 2016



Outline

- 1 Motivation
- 2 Problem Statement
- 3 Concepts & Methods
- 4 Experiment Results
- 5 Conclusion

Motivation

Problem Statement

Parallel kNN Search

Brute Force

- simply implementable
- highly parallelisable
- low memory requirements
- quadratic runtime

kd-Tree

- hard to implement
- only search part parallelisable
- memory-intensive
- linear runtime

Matrix Multiplication

image of matrix multiplication

Pseudocode

```
for all points in data do  
    calculate distances to all other points  
    repeat  
        count neighbours in distance of  $\varepsilon$   
        adapt  $\varepsilon$   
    until number of neighbours in distance  $\varepsilon = k$   
    get neighbours in distance of  $\varepsilon$   
end for
```

kd-Tree

Results

Conclusion